

Pipeline Operations

CS 4620 Lecture II

Pipeline

you are here →

APPLICATION

COMMAND STREAM

3D transformations; shading →

VERTEX PROCESSING

TRANSFORMED GEOMETRY

conversion of primitives to pixels →

RASTERIZATION

FRAGMENTS

blending, compositing, shading →

FRAGMENT PROCESSING

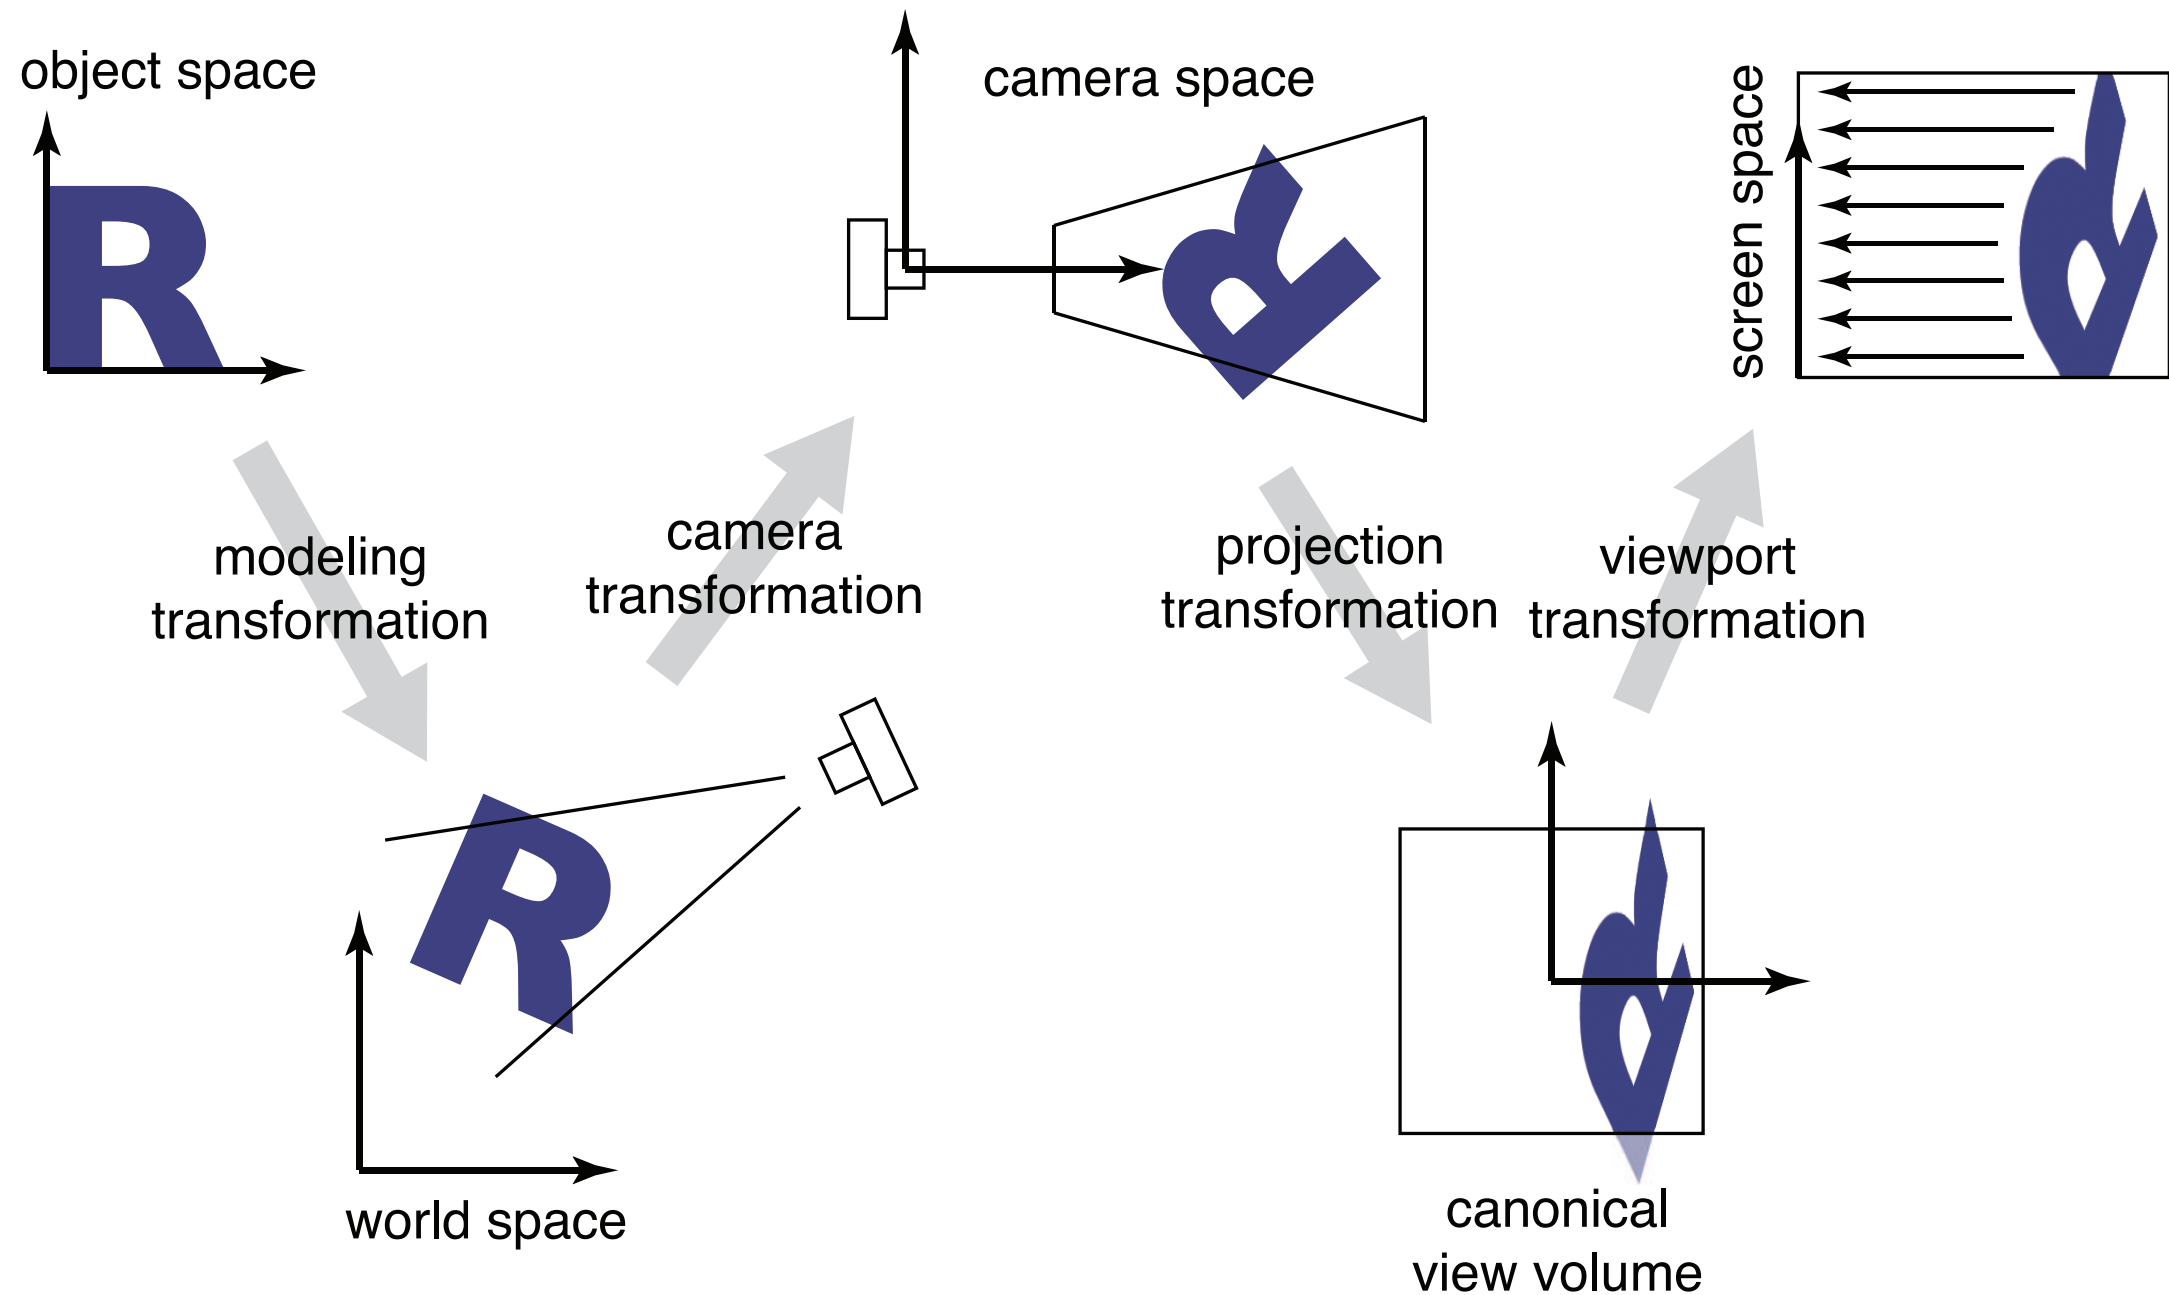
FRAMEBUFFER IMAGE

user sees this →

DISPLAY

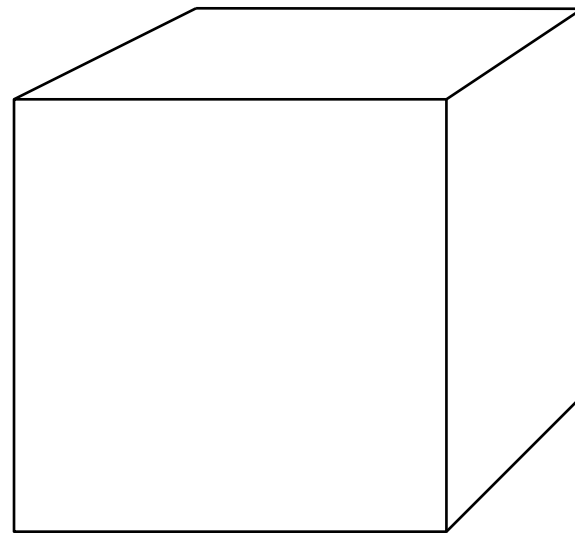
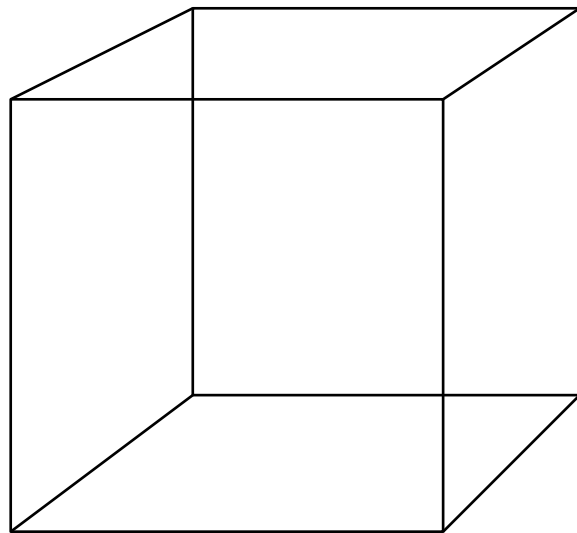
Pipeline of transformations

- **Standard sequence of transforms**



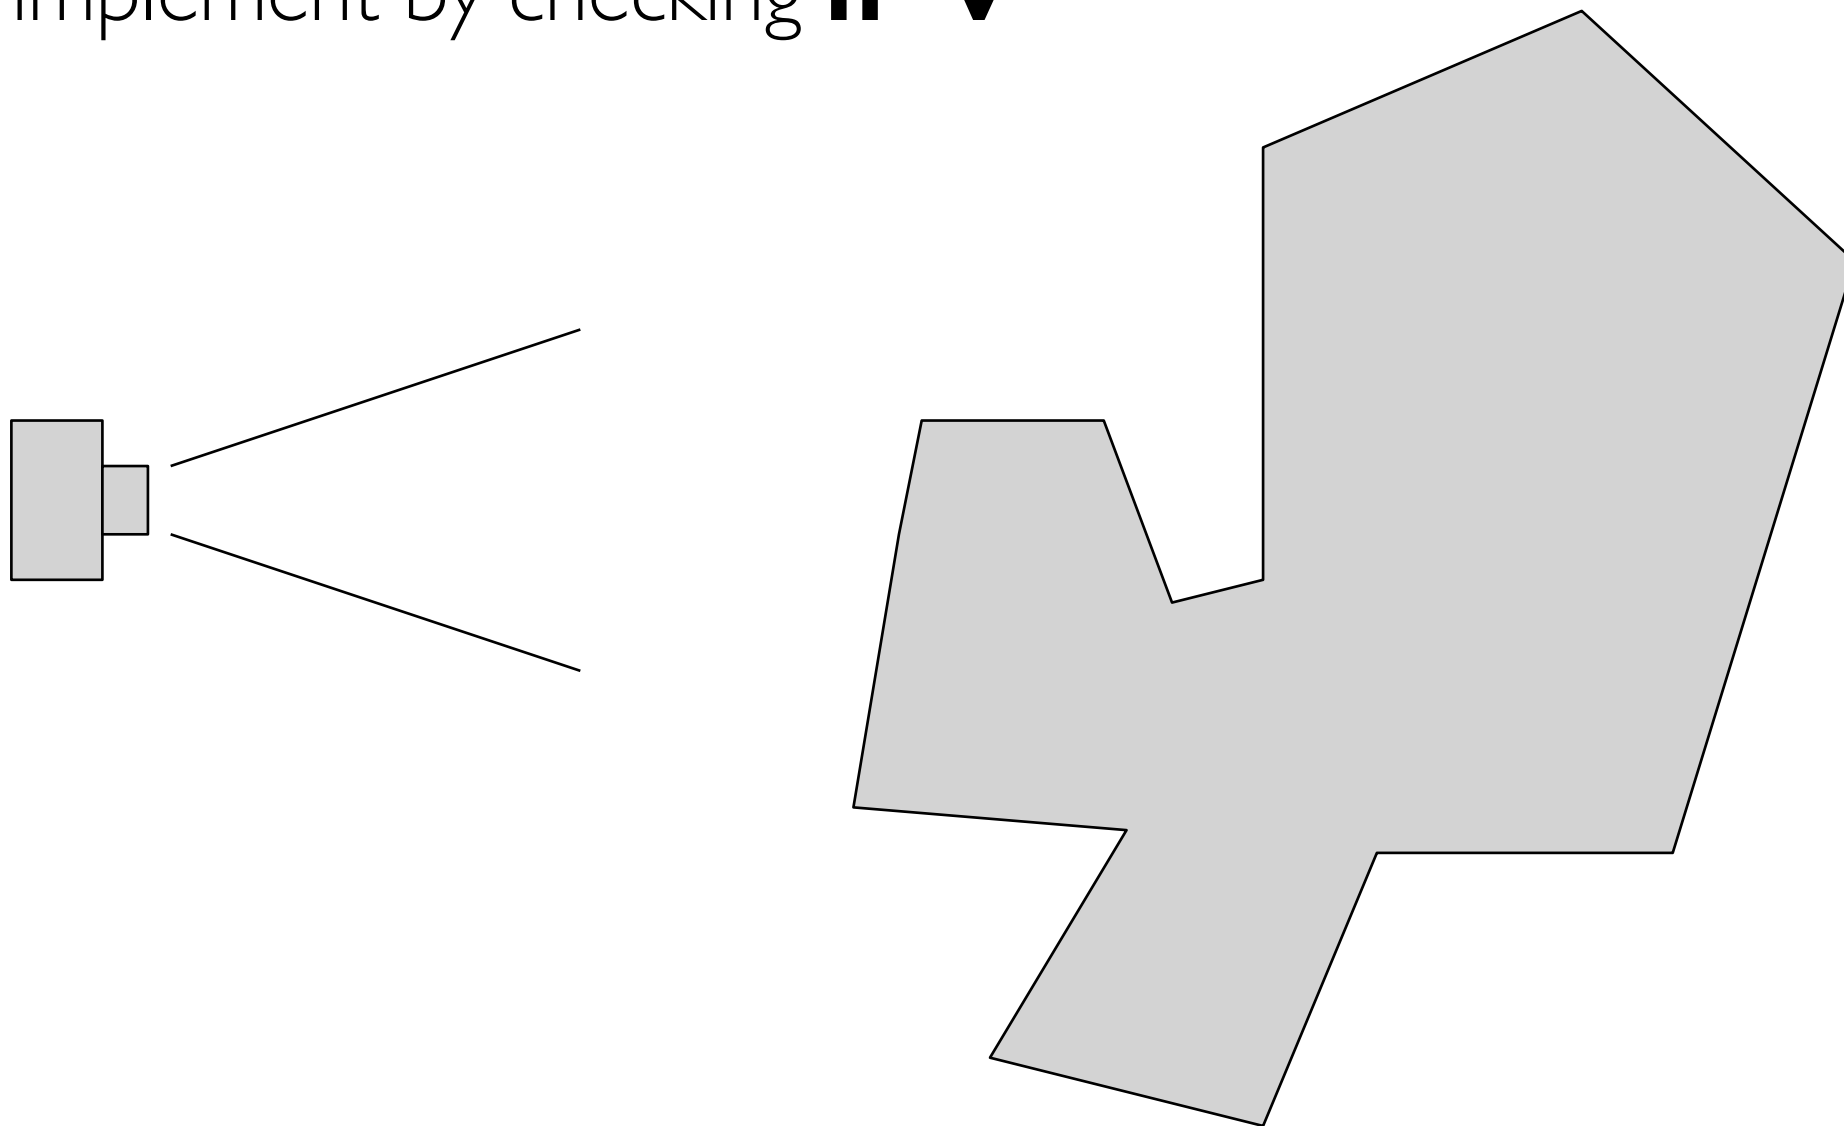
Hidden surface elimination

- **We have discussed how to map primitives to image space**
 - projection and perspective are depth cues
 - occlusion is another very important cue



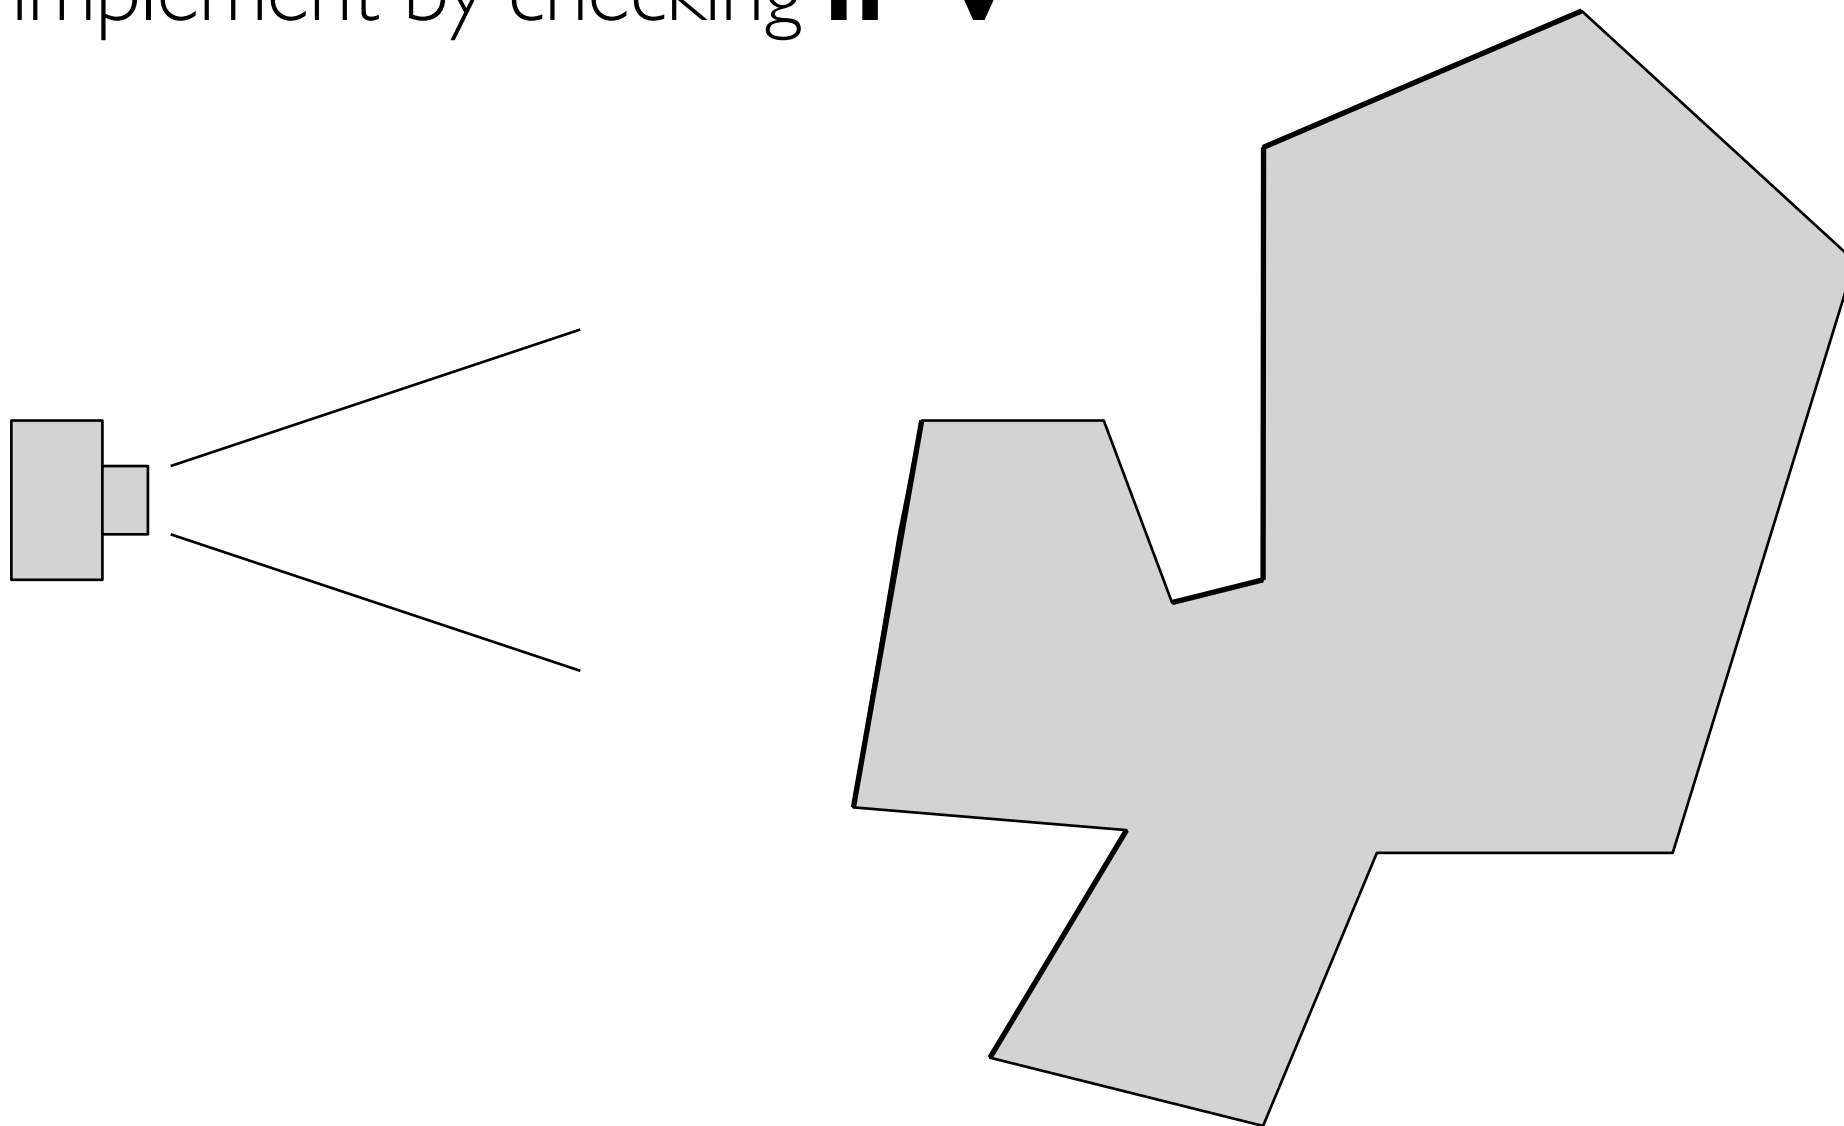
Back face culling

- **For closed shapes you will never see the inside**
 - therefore only draw surfaces that face the camera
 - implement by checking **$\mathbf{n} \cdot \mathbf{v}$**




Back face culling

- **For closed shapes you will never see the inside**
 - therefore only draw surfaces that face the camera
 - implement by checking **$\mathbf{n} \cdot \mathbf{v}$**



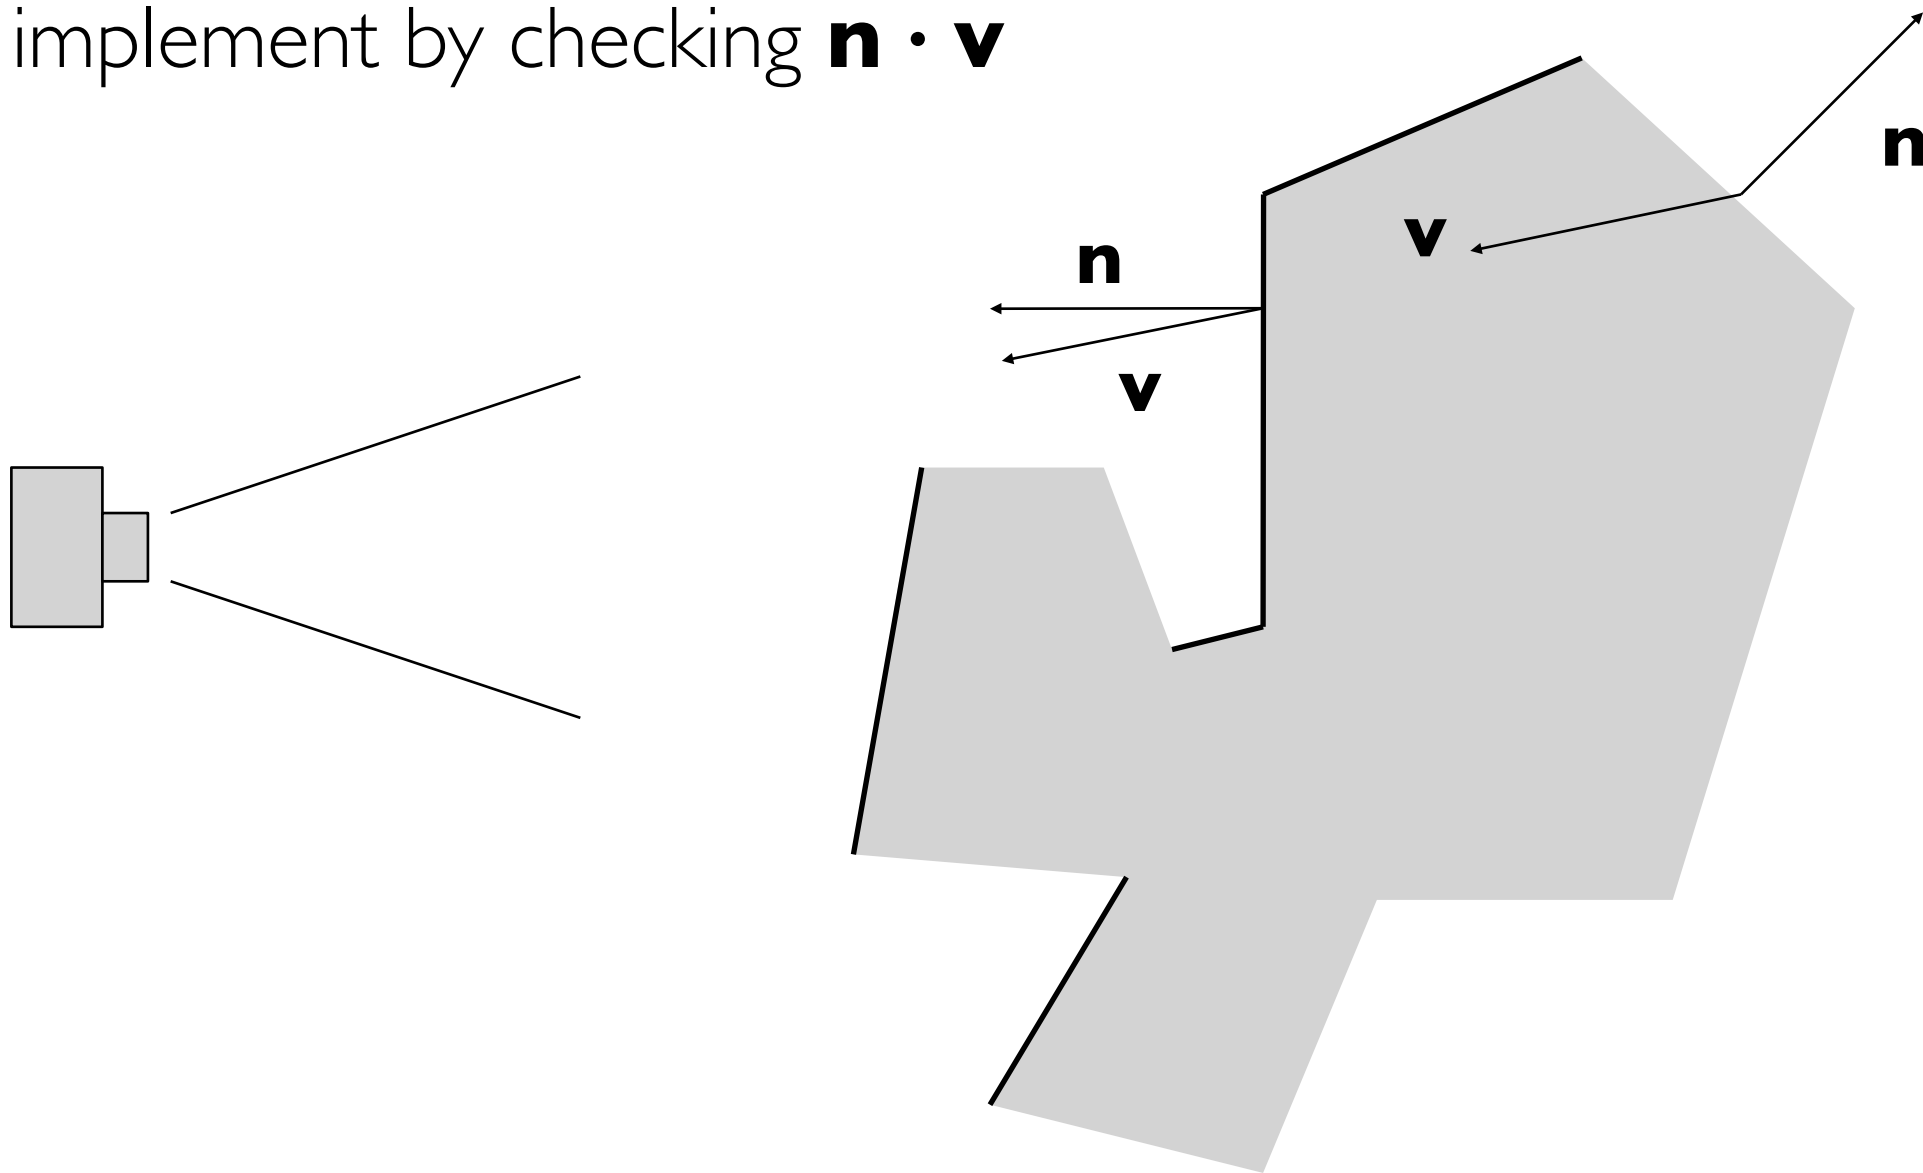
Back face culling

- **For closed shapes you will never see the inside**
 - therefore only draw surfaces that face the camera
 - implement by checking **$\mathbf{n} \cdot \mathbf{v}$**



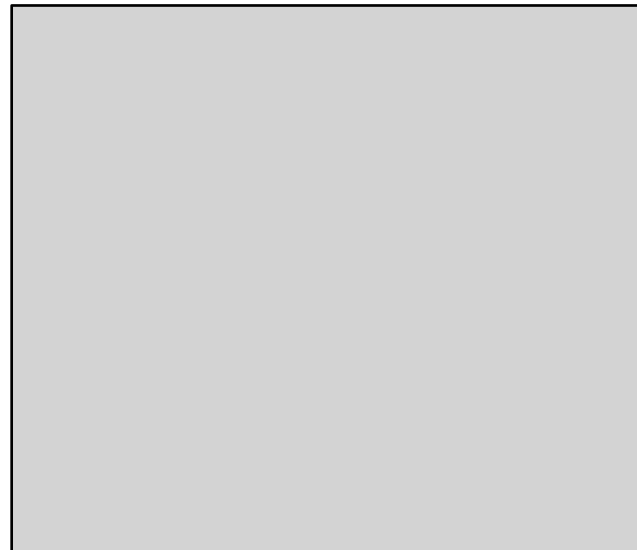
Back face culling

- **For closed shapes you will never see the inside**
 - therefore only draw surfaces that face the camera
 - implement by checking $\mathbf{n} \cdot \mathbf{v}$



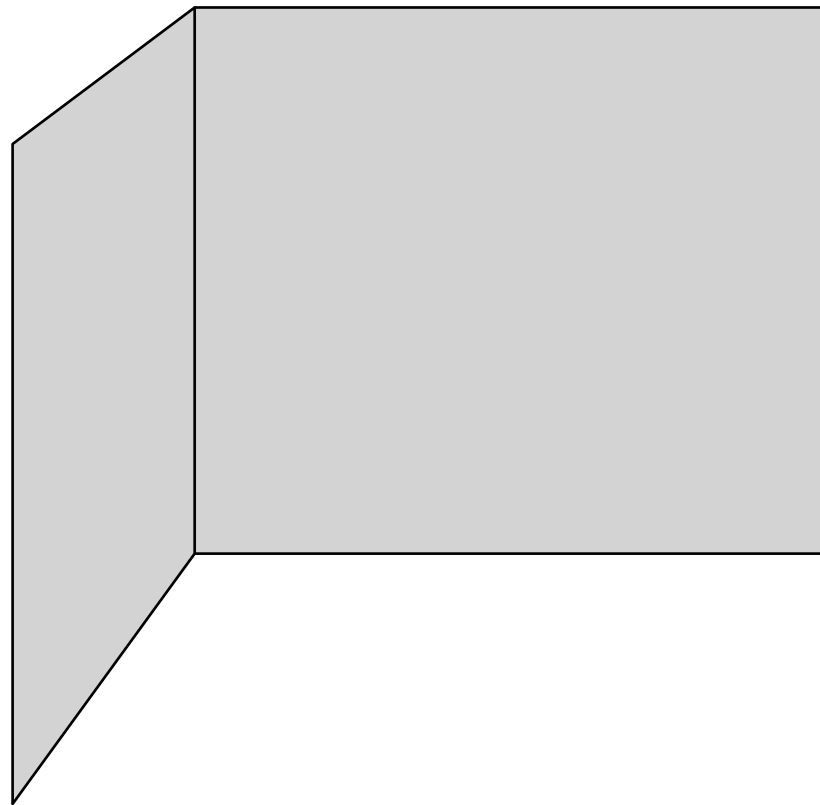
Painter's algorithm

- **Simplest way to do hidden surfaces**
- **Draw from back to front, use overwriting in framebuffer**



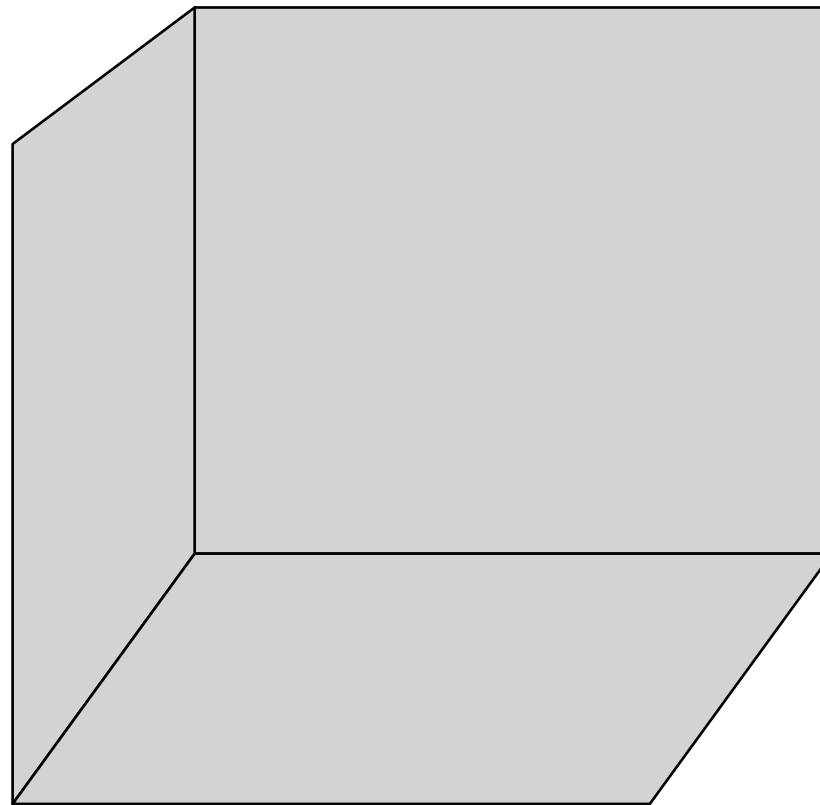
Painter's algorithm

- **Simplest way to do hidden surfaces**
- **Draw from back to front, use overwriting in framebuffer**



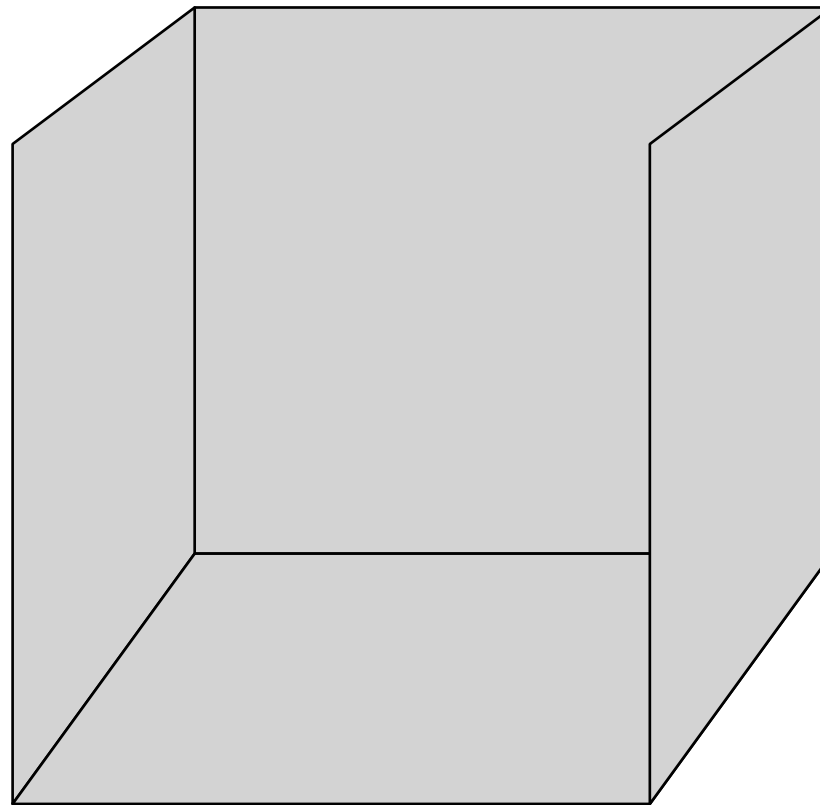
Painter's algorithm

- **Simplest way to do hidden surfaces**
- **Draw from back to front, use overwriting in framebuffer**



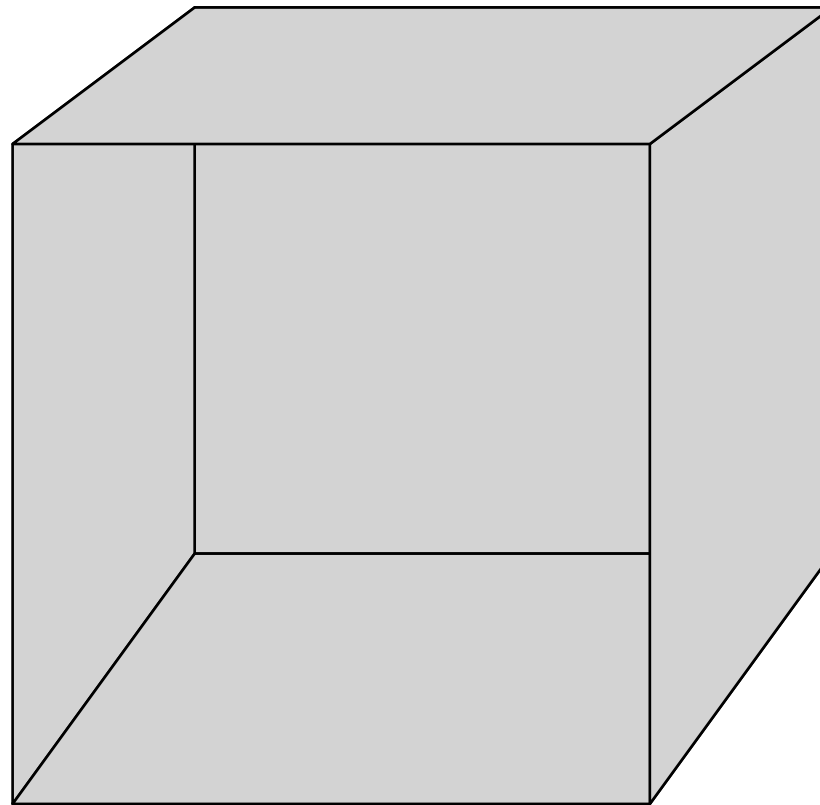
Painter's algorithm

- **Simplest way to do hidden surfaces**
- **Draw from back to front, use overwriting in framebuffer**



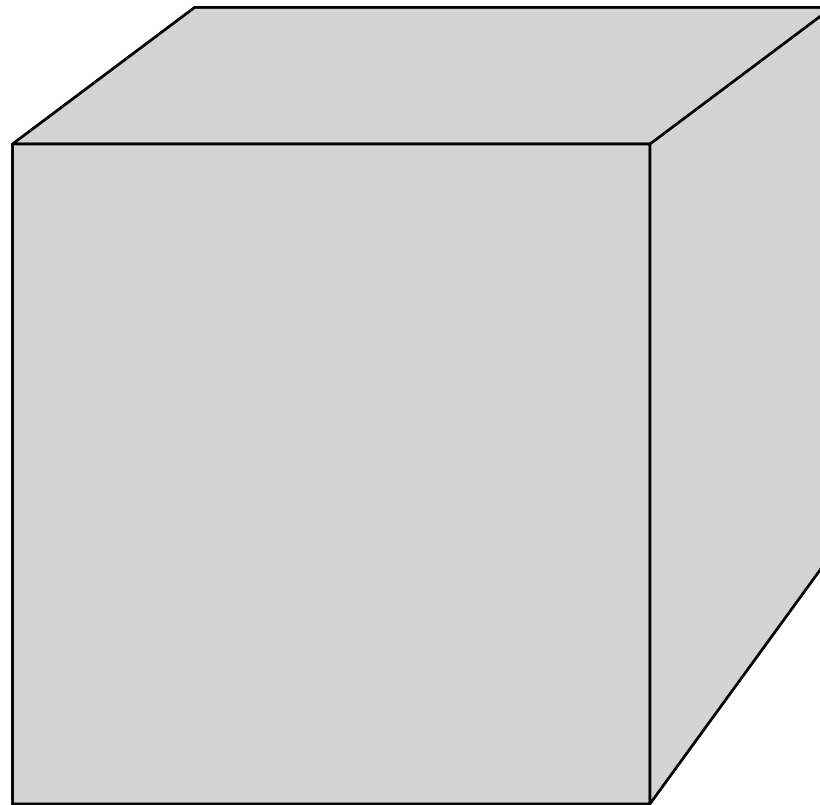
Painter's algorithm

- **Simplest way to do hidden surfaces**
- **Draw from back to front, use overwriting in framebuffer**



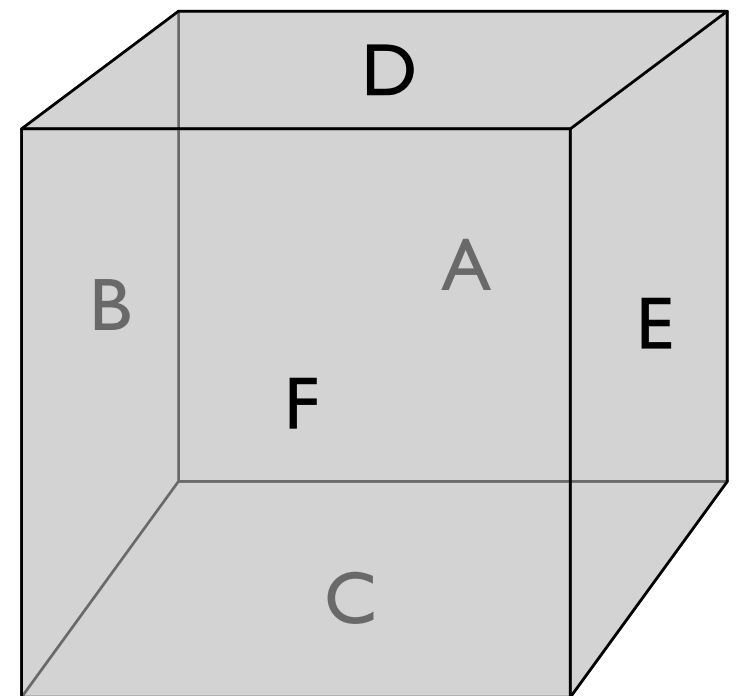

Painter's algorithm

- **Simplest way to do hidden surfaces**
- **Draw from back to front, use overwriting in framebuffer**





Painter's algorithm

- **Amounts to a topological sort of the graph of occlusions**
 - that is, an edge from A to B means A sometimes occludes B
 - any sort is valid
 - ABCDEF
 - BADCFE
 - if there are cycles there is no sort



Painter's algorithm


- **Amounts to a topological sort of the graph of occlusions**
 - that is, an edge from A to B means A sometimes occludes B
 - any sort is valid
 - ABCDEF
 - BADCFE
 - if there are cycles there is no sort



[Foley et al.]

Painter's algorithm

- **Useful when a valid order is easy to come by**
- **Compatible with alpha blending**




[Foley et al.]

The **z** buffer

- **In many (most) applications maintaining a z sort is too expensive**
 - changes all the time as the view changes
 - many data structures exist, but complex
- **Solution: draw in any order, keep track of closest**
 - allocate extra channel per pixel to keep track of closest depth so far
 - when drawing, compare object's depth to current closest depth and discard if greater
 - this works just like any other compositing operation

The **z** buffer




[Foley et al.]

- another example of a memory-intensive brute force approach that works and has become the standard

Precision in **z** buffer


- **The precision is distributed between the near and far clipping planes**
 - this is why these planes have to exist
 - also why you can't always just set them to very small and very large distances
- **Generally use z' (not world z) in z buffer**

Interpolating in projection




linear interp. in screen space \neq linear interp. in world (eye) space

Interpolating in projection




linear interp. in screen space \neq linear interp. in world (eye) space

Interpolating in projection




linear interp. in screen space \neq linear interp. in world (eye) space

Interpolating in projection




linear interp. in screen space \neq linear interp. in world (eye) space

Interpolating in projection




linear interp. in screen space \neq linear interp. in world (eye) space

Interpolating in projection



linear interp. in screen space \neq linear interp. in world (eye) space

Interpolating in projection




linear interp. in screen space \neq linear interp. in world (eye) space

Pipeline for minimal operation

- **Vertex stage (input: position / vtx)**
 - transform position (object to screen space)
- **Rasterizer**
 - nothing (extra) to interpolate
- **Fragment stage (output: color)**
 - write a fixed color to color planes
 - (color is a “uniform” quantity that is infrequently updated)


Result of minimal pipeline



Pipeline for basic z buffer

- **Vertex stage (input: position / vtx)**
 - transform position (object to screen space)
- **Rasterizer**
 - interpolated parameter: z' (screen z)
- **Fragment stage (output: color, z')**
 - write fixed color to color planes only if interpolated $z' <$ current z'

Result of **z**-buffer pipeline



Flat shading

- **Shade using the real normal of the triangle**
 - same result as ray tracing a bunch of triangles
- **Leads to constant shading and faceted appearance**
 - truest view of the mesh geometry





Plate II.29 *Shutterbug*. Individually shaded polygons with diffuse reflection (Sections 14.4.2 and 16.2.3). (Copyright © 1990, Pixar. Rendered by Thomas Williams and H.B. Siegel using Pixar's PhotoRealistic RenderMan™ software.)

[Foley et al.]

Pipeline for flat shading

- **Vertex stage (input: position and normal / vtx)**
 - transform position and normal (object to eye space)
 - compute shaded color per triangle using normal
 - transform position (eye to screen space)
- **Rasterizer**
 - interpolated parameters: \mathbf{z}' (screen \mathbf{z})
 - pass through color — non-default rasterization mode these days (rasterizer will use the color from the first vtx. for all fragments)
- **Fragment stage (output: color, \mathbf{z}')**
 - write fixed color to color planes only if interpolated $\mathbf{z}' <$ current \mathbf{z}'


Result of flat-shading pipeline



Transforming normal vectors

- **Transforming surface normals**

- differences of points (and therefore tangents) transform OK
- normals do not --> use inverse transpose matrix



have: $\mathbf{t} \cdot \mathbf{n} = \mathbf{t}^T \mathbf{n} = 0$

want: $M\mathbf{t} \cdot X\mathbf{n} = \mathbf{t}^T M^T X\mathbf{n} = 0$



so set $X = (M^T)^{-1}$

then: $M\mathbf{t} \cdot X\mathbf{n} = \mathbf{t}^T M^T (M^T)^{-1} \mathbf{n} = \mathbf{t}^T \mathbf{n} = 0$

Transforming normal vectors

- **Transforming surface normals**

- differences of points (and therefore tangents) transform OK
- normals do not --> use inverse transpose matrix



have: $\mathbf{t} \cdot \mathbf{n} = \mathbf{t}^T \mathbf{n} = 0$



want: $M\mathbf{t} \cdot X\mathbf{n} = \mathbf{t}^T M^T X\mathbf{n} = 0$

so set $X = (M^T)^{-1}$

then: $M\mathbf{t} \cdot X\mathbf{n} = \mathbf{t}^T M^T (M^T)^{-1} \mathbf{n} = \mathbf{t}^T \mathbf{n} = 0$

Gouraud shading

- **Often we're trying to draw smooth surfaces, so facets are an artifact**
 - compute colors at vertices using vertex normals
 - interpolate colors across triangles
 - “Gouraud shading”
 - “Smooth shading”




[Gouraud thesis]

Gouraud shading

- **Often we're trying to draw smooth surfaces, so facets are an artifact**
 - compute colors at vertices using vertex normals
 - interpolate colors across triangles
 - “Gouraud shading”
 - “Smooth shading”

Plate II.30 *Shutterbug*. Gouraud shaded polygons with diffuse reflection (Sections 14.4.3 and 16.2.4). (Copyright © 1990, Pixar. Rendered by Thomas Williams and H.B. Siegel using Pixar's PhotoRealistic RenderMan™ software.)




[Foley et al.]

Pipeline for Gouraud shading


- **Vertex stage (input: position and normal / vtx)**
 - transform position and normal (object to eye space)
 - compute shaded color per vertex (using fixed diffuse color)
 - transform position (eye to screen space)
- **Rasterizer**
 - interpolated parameters: \mathbf{z}' (screen \mathbf{z}); r, g, b color
- **Fragment stage (output: color, \mathbf{z}')**
 - write to color planes only if interpolated $\mathbf{z}' <$ current \mathbf{z}'

Result of Gouraud shading pipeline



Local vs. infinite viewer, light

- **Phong illumination requires geometric information:**
 - light vector (function of position)
 - eye vector (function of position)
 - surface normal (from application)
- **Light and eye vectors change**
 - need to be computed (and normalized) for each vertex




Local vs. infinite viewer, light

- **Look at case when eye or light is far away:**
 - distant light source: nearly parallel illumination
 - distant eye point: nearly orthographic projection
 - in both cases, eye or light vector changes very little
- **Optimization: approximate eye and/or light as infinitely far away**


Directional light

- **Directional (infinitely distant) light source**
 - light vector always points in the same direction
 - often specified by position $[x \ y \ z \ 0]$
 - many pipelines are faster if you use directional lights





Directional light

- **Directional (infinitely distant) light source**
 - light vector always points in the same direction
 - often specified by position $[x \ y \ z \ 0]$
 - many pipelines are faster if you use directional lights



Infinite viewer


- **Orthographic camera**
 - projection direction is constant
- **“Infinite viewer”**
 - even with perspective, can approximate eye vector using the image plane normal
 - can produce weirdness for wide-angle views
 - light, eye, half vectors all constant! 



Vertex normals

- **Need normals at vertices to compute Gouraud shading**
- **Best to get vtx. normals from the underlying geometry**
 - e. g. spheres example
- **Otherwise have to infer vtx. normals from triangles**
 - simple scheme: average surrounding face normals


$$N_v = \frac{\sum_i N_i}{\|\sum_i N_i\|}$$



[Foley et al.]

Non-diffuse Gouraud shading


- **Can apply Gouraud shading to any illumination model**
 - it's just an interpolation method
- **Results are not so good with fast-varying models like specular ones**
 - problems with any highlights smaller than a triangle



[Foley et al.]

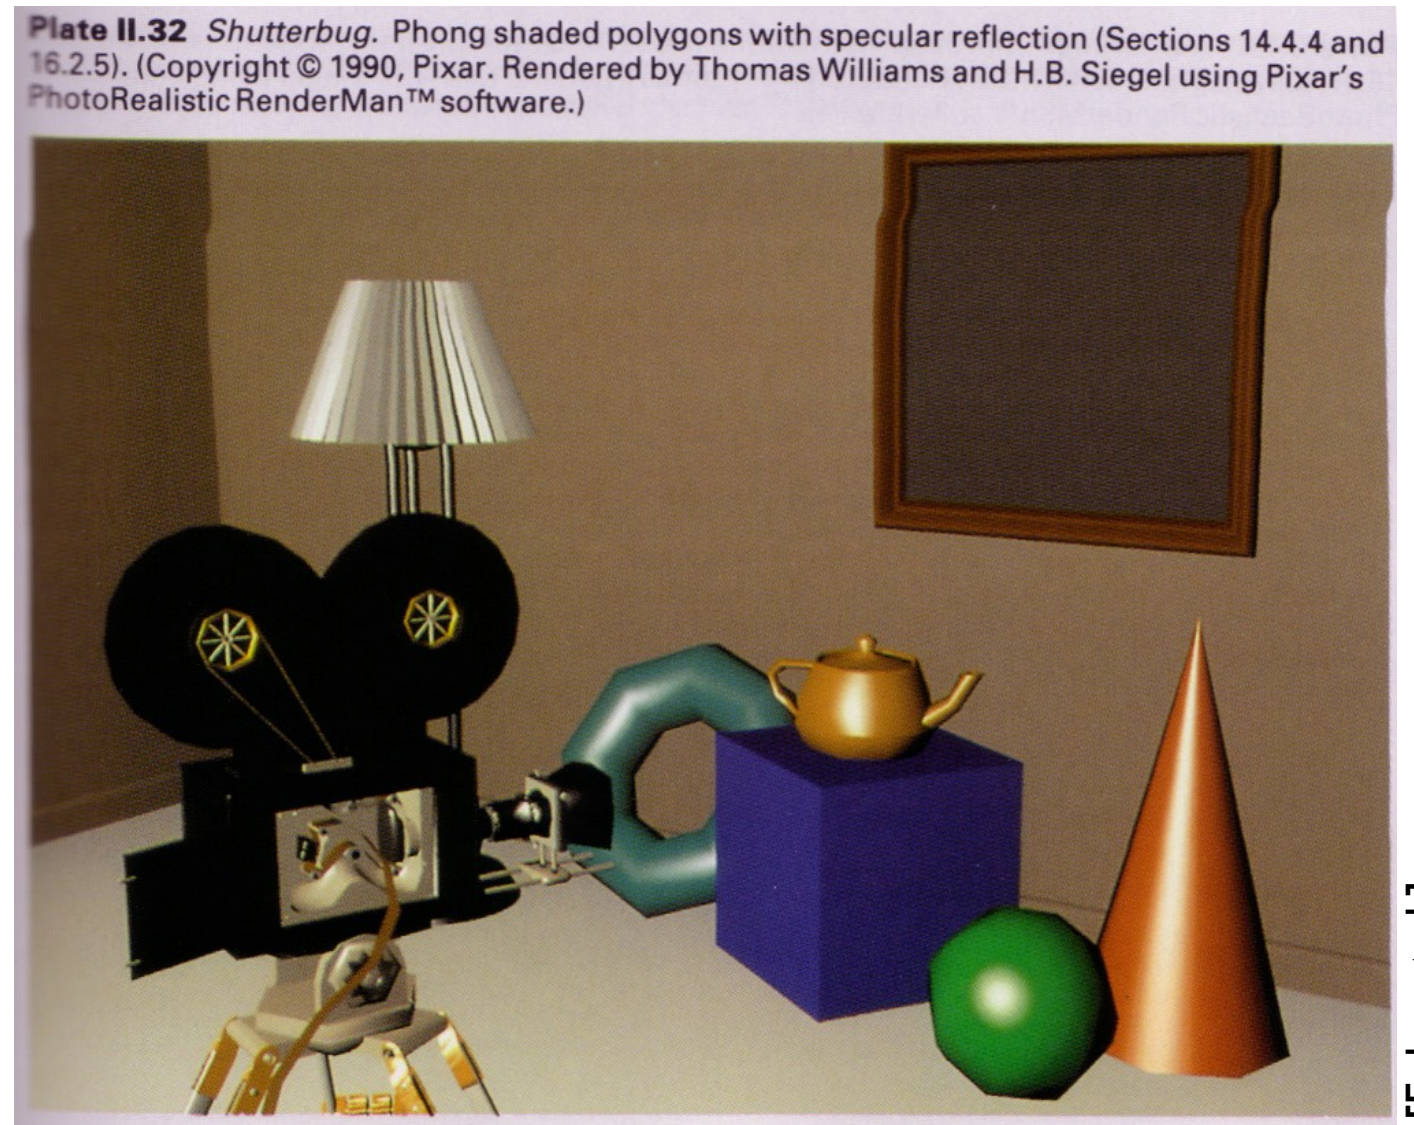

Per-pixel (Phong) shading

- **Get higher quality by interpolating the normal**
 - just as easy as interpolating the color
 - but now we are evaluating the illumination model per pixel rather than per vertex (and normalizing the normal first)
 - in pipeline, this means we are moving illumination from the vertex processing stage to the fragment processing stage



Per-pixel (Phong) shading

- **Bottom line: produces much better highlights**




[Foley et al.]

Pipeline for per-pixel shading

- **Vertex stage (input: position and normal / vtx)**
 - transform position and normal (object to eye space)
 - transform position (eye to screen space)
- **Rasterizer**
 - interpolated parameters: \mathbf{z}' (screen \mathbf{z}); \mathbf{x} , \mathbf{y} , \mathbf{z} normal
- **Fragment stage (output: color, \mathbf{z}')**
 - compute shading using fixed color and interpolated normal
 - write to color planes only if interpolated $\mathbf{z}' <$ current \mathbf{z}'

Result of per-pixel shading pipeline



Programming hardware pipelines

- **Modern hardware graphics pipelines are flexible**
 - programmer defines exactly what happens at each stage
 - do this by writing *shader programs* in domain-specific languages called *shading languages*
 - rasterization is fixed-function, as are some other operations (depth test, many data conversions, ...)
- **One example: OpenGL and GLSL (GL Shading Language)**
 - several types of shaders process primitives and vertices; most basic is the *vertex program*
 - after rasterization, fragments are processed by a *fragment program*

GLSL Shaders

